


Lecture No.3: Matrix Multiplication

Muhammad Osama Mahmoud, TA

1

Basic Matrix Multiplication



Basic Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(int m, int n, int k, float* A, float* B, float* C)
{
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
    int Col = blockIdx.x * blockDim.x + threadIdx.x;
    if ((Row < m) && (Col < k)) {
        float C_Element = 0.0;
        for (int i = 0; i < n; i++)
            C_Element += A[Row*n + i] * B[Col + i*k];
        C[Row*k + Col] = C_Element;
    }
}
```

Basic Kernel Performance

- All threads access global memory for every matrix input even if reused many times
- Assume two matrices (A, B) have the same number of elements ($n \times n$)
- For the first row in A, the number of multiply-add operations is n
- The number of repeated global memory accesses are n^2
- These repeated global accesses can be reduced from n^2 to just n if we take advantage of the tiling strategy
- Now we ready for the tiled matrix multiplication algorithm

Tiled Matrix Multiplication


- Load a submatrix from A and B into shared memory
- Process data reside in shared memory
- Write the result submatrix into global memory
- Ex. Accessing t block/tile in 2D indexing:

$A[\text{Row}][t * \text{BLOCK_SIZE} + tx]$

$B[t * \text{BLOCK_SIZE} + ty][\text{Col}]$

Where $tx = \text{threadIdx.x}$

$ty = \text{threadIdx.y}$



Tiled Matrix Multiplication (Cont.)

- Mapping from 2D to 1D
 - $A[\text{Row}][t * \text{BLOCK_SIZE} + tx]$



$A[\text{Row} * n + t * \text{BLOCK_SIZE} + tx]$, where n is the number of elements in a row of A

- $B[t * \text{BLOCK_SIZE} + ty][\text{Col}]$



$B[(t * \text{BLOCK_SIZE} + ty) * k + \text{Col}]$, where k is the number of elements in a column of B

Tiled Matrix Multiplication Kernel

```
__global__ void TiledMatrixMulKernel(int m, int n, int k, float* A, float* B, float* C) {  
    __shared__ float ds_A[TILE_WIDTH][TILE_WIDTH];  
    __shared__ float ds_B[TILE_WIDTH][TILE_WIDTH];  
    int bx = blockIdx.x; int by = blockIdx.y;  
    int tx = threadIdx.x; int ty = threadIdx.y;  
    int Row = by * blockDim.y + ty;  
    int Col = bx * blockDim.x + tx;  
    float C_Element = 0.0;  
    // rest of kernel in next slide
```

Tiled Matrix Multiplication Kernel (cont.)

```
// Loop over the A and B tiles required to compute the C element
```

```
for (int t = 0; t < n/BLOCK_SIZE; t++) {
```

```
    // load of A and B tiles into shared memory
```

```
    ds_A[ty][tx] = A[Row*n + t*BLOCK_SIZE+tx];
```

```
    ds_B[ty][tx] = B[(t*BLOCK_SIZE+ty)*k + Col];
```

```
    __syncthreads();
```

```
    for (int i = 0; i < BLOCK_SIZE; ++i)
```

```
        C_Element += ds_A[ty][i] * ds_B[i][tx];
```

```
    __syncthreads();
```

```
}
```

```
    C[Row*k+Col] = C_Element;
```

```
} // kernel end
```


References

- [1] Wen-mei W. Hwu, “Heterogeneous Parallel Programming”. Online course, 2014.
Available: <https://class.coursera.org/hetero-002>
- [2] NVIDIA, “CUDA C Programming Guide”, June 2014.