

Lecture No.4: Tree-based Reduction

Muhammad Osama Mahmoud, TA


1

What is Reduction ?

- Reduce a set of input data values into a single value using some reduction operation
 - Sum
 - Min
 - Max
 - Product
- May be applied to dot-product as we see later
- The sequential version has a complexity of $O(N)$, recursive over N input data and perform the reduction operation N times

Tree-based reduction algorithm

- Assume a sum operation



Basic Parallel Reduction

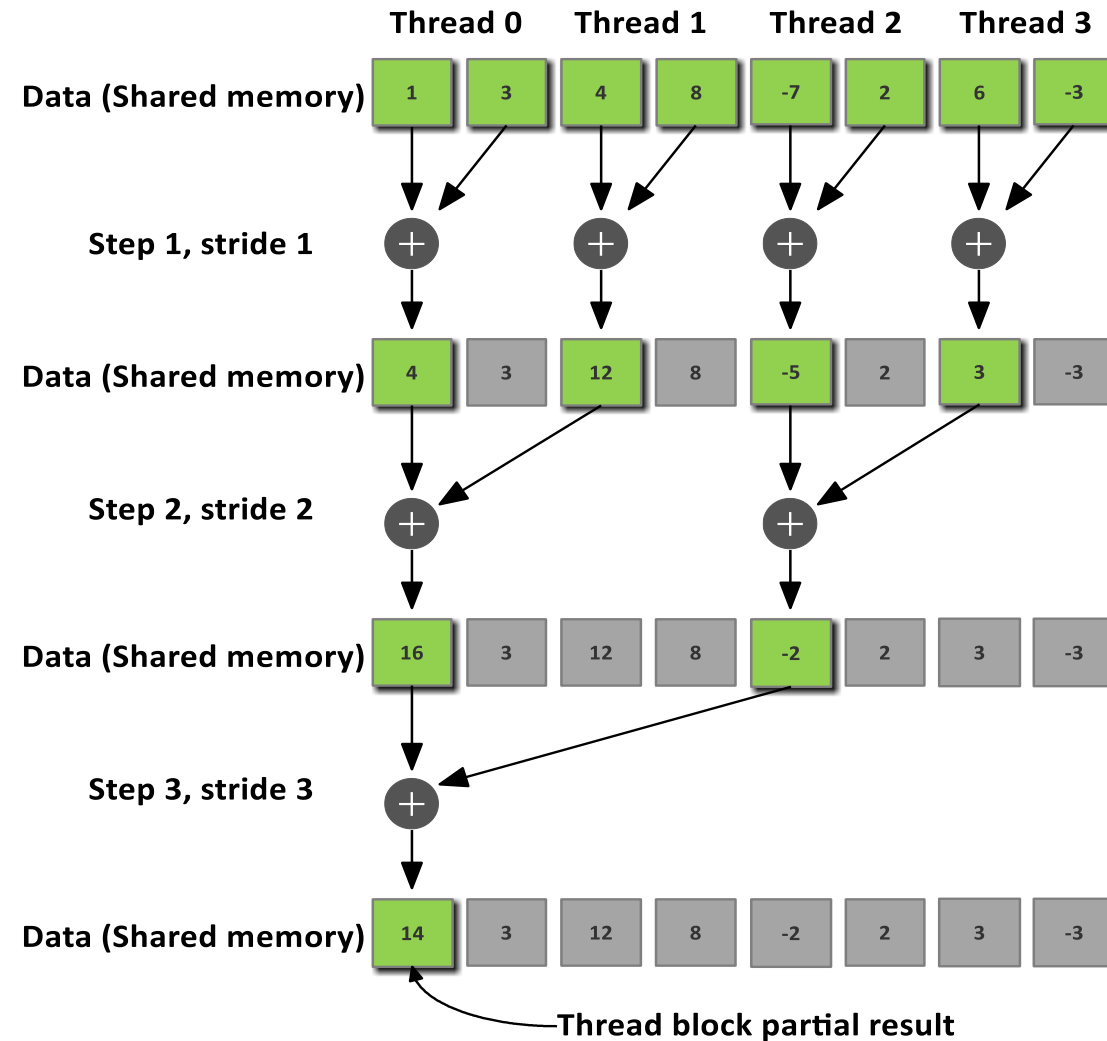
- Assume the sum operation
- Uses the interleaved addressing
- Each thread adds two consecutive elements within the shared memory
- Recursively halve the number of threads after each step
- Repeat the reduction operation on the resulting data
- Add the final partial results on the host side

Basic Reduction Kernel

```
__global__ void reduce_v1(int * inp_data, int * outp_data) {  
    extern __shared__ int sh_data[ ]; //dynamically locate shared memory at kernel launch  
    unsigned int tx = threadIdx.x;  
    unsigned int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    sh_data[tx] = inp_data[idx];  
    __syncthreads();  
    // do reduction in shared memory  
    for(unsigned int stride = 1; stride < blockDim.x; stride <= 1) {  
        int index = 2 * stride * tx;  
        if ((index) < blockDim.x) {  
            sh_data[index] += sh_data[index + stride];  
        }  
        __syncthreads();  
    }  
    // make thread 0 write result for each block to global memory  
    if (tx == 0) outp_data[blockIdx.x] = sh_data[tx];  
}
```

Basic Reduction Kernel Execution

- Input data of 8 values
- Single thread block
- Block = 8 threads
- Shared memory =
8 threads * `sizeof(int)`



Basic Reduction Kernel Execution (Cont.)

- Eliminates the problem of thread-divergence while reducing the number of threads at each stride
- The interleaved addressing has a disadvantage of shared memory bank conflicts
- A bank conflict arises when a halfwarp tries to load/store from/to the same memory bank (the access will be serialized)
- Solution, use the sequential addressing instead to avoid any bank conflicts (each thread in halfwarp access consecutive 32-bit words)
- A barrier synchronization required to make sure each value in shared memory is updated, that's why `__syncthreads()` is used

Reduction Kernel with sequential addressing

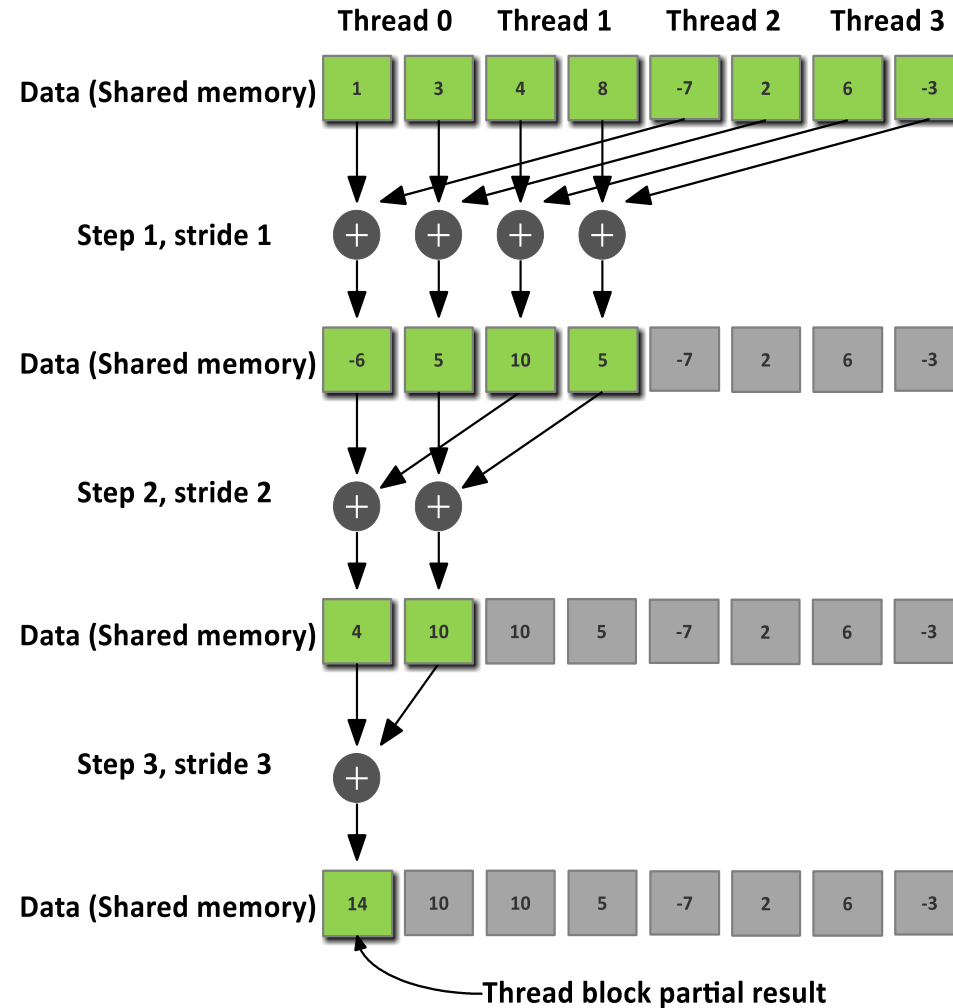
- Just replace the reduction step in the basic kernel

```
for (unsigned int stride = 1; stride < blockDim.x; stride  
<=< 1) {  
    int index = 2 * stride * tx;  
    if ((index) < blockDim.x) {  
        sh_data[index] += sh_data[index + stride];  
    }  
    __syncthreads();  
}
```

With a reversed loop and threadID-based indexing:

```
for (unsigned int stride = blockDim.x/2; stride > 0;  
stride >>= 1) {  
    if (tx < stride) {  
        sh_data[tx] += sh_data[tx + stride];  
    }  
    __syncthreads();  
}
```


Reduction Kernel with sequential addressing Execution



References

- [1] Wen-mei W. Hwu, “Heterogeneous Parallel Programming”. Online course, 2014.
Available: <https://class.coursera.org/hetero-002>
- [2] M. Harris, “Optimizing Parallel Reduction in CUDA”, Oct. 2007.