

# Lecture No.5: Advanced Parallel Reduction

Muhammad Osama Mahmoud, TA

1

# Advanced Parallel Reduction #1

- The number of arithmetic instructions compared to loop and addressing instructions is low
- This instruction overhead can be evaded using loop unrolling
- A full loop unrolling on the reduction loop is considered
- After each reduction step the number of idle threads increases
- To minimize thread idle time, make each thread adds two elements from two different blocks and load to shared memory

# Advanced Parallel Reduction #1 (Cont.)

```
__global__ void reduce_v3(int * inp_data, int * outp_data) {  
    extern __shared__ int sh_data[ ]; //dynamically locate shared memory at kernel launch  
    unsigned int tx = threadIdx.x;  
    unsigned int idx = blockIdx.x * (blockDim.x * 2) + threadIdx.x;  
    sh_data[tx] = inp_data[idx] + inp_data[idx + blockDim.x]; // perform first add step then load to shared memory  
    __syncthreads();  
    // perform full loop unrolling  
    if (blockDim.x >= 512 && tx < 256) {  
        if (tx < 256) { sh_data[tx] += sh_data[tx + 256]; } __syncthreads();  
    }  
    if (blockDim.x >= 256 && tx < 128) {  
        if (tx < 128) { sh_data[tx] += sh_data[tx + 128]; } __syncthreads();  
    }  
    if (blockDim.x >= 128 && tx < 64) {  
        if (tx < 64) { sh_data[tx] += sh_data[tx + 64]; } __syncthreads();  
    }  
}
```

// rest of code in the next slide

# Advanced Parallel Reduction #1 (Cont.)

```
// perform loop unrolling on the last warp
```

```
if (tid < 32) {
```

```
    if (blockDim.x >= 64 && tx < 32) sh_data[tx] += sh_data[tx + 32];
```

```
    if (blockDim.x >= 32 && tx < 16) sh_data[tx] += sh_data[tx + 16];
```

```
    if (blockDim.x >= 16 && tx < 8) sh_data[tx] += sh_data[tx + 8];
```

```
    if (blockDim.x >= 8 && tx < 4) sh_data[tx] += sh_data[tx + 4];
```

```
    if (blockDim.x >= 4 && tx < 2) sh_data[tx] += sh_data[tx + 2];
```

```
    if (blockDim.x >= 2 && tx < 1) sh_data[tx] += sh_data[tx + 1];
```

```
}
```

```
if (tx == 0) outp_data[blockIdx.x] = sh_data[tx];
```

```
} // kernel end
```

# Parallel Reduction Complexity

- $\log(N)$  sequential strides or steps, each stride performs  $2^{\text{stride}}$  operations, that is,  $N$  independent operations
- With  $P$  threads physically initiated in parallel, time complexity is  $O(N/P + \log N)$
- If  $P = N$  then complexity is  $O(1 + \log N) = O(\log N)$
- The performance speed-up will be  $s = \frac{T_{\text{serial}}}{T_{\text{parallel}}} = \frac{N}{\log N}$  times faster than the sequential algorithm

# Parallel Reduction Cost

- Cost of parallel algorithm = # of processors (threads) x time complexity
- In case of parallel reduction, cost =  $N \times O(\log N) = O(N \log N)$ , not cost efficient
- Brent's theorem states that each processor should do  $O(\log N)$  sequential operations
- If applied to the reduction kernel, the number of threads needed is  $O(N/\log N)$ , and the cost will be  $O(N/\log N) \times O(\log N) = O(N)$ , which is cost efficient

# Advanced Parallel Reduction #2

- To apply Brent's theorem, we replace and modify the following code block from the last kernel

```
sh_data[tx] = inp_data[idx] + inp_data[idx +  
blockDim.x];
```

with a loop to add and load as many elements as possible

```
unsigned int gridSize = + blockDim.x * 2 * gridDim.x;  
sdata[tx] = 0;  
  
while (i < n) {  
    sh_data[tx] += inp_data[idx] + inp_data [idx + blockDim.x ];  
    idx += gridSize;  
}  
__syncthreads();
```

# References

- [1] Wen-mei W. Hwu, “Heterogeneous Parallel Programming”. Online course, 2014. Available: <https://class.coursera.org/hetero-002>
- [2] M. Harris, “Optimizing Parallel Reduction in CUDA”, Oct. 2007.